Markus Gruber, M.Sc.

Senior Machine Learning Engineer
Data Scientist

AI Infrastructure · RAG Pipelines · Scientific Computing Specialist

3100 St. Pölten, Austria · +43 650 5871145 · Markus.Gruber4@gmx.net · LinkedIn · GitHub

Professional Profile



Analytical Senior Machine Learning Engineer & Data Scientist holding an M.Sc. in Scientific Computing with a strong mathematical foundation. Over 9 years of industry experience deploying production-grade AI pipelines and data architectures.

Expertise bridges mathematical algorithm design, GPU edge acceleration (CUDA), and cloud-native RAG infrastructure (Vector DBs, PyTorch). Highly effective at productionizing complex ML models with high execution throughput.

Engineering Profile

TECHNICAL FOUNDATION & AI / DATA DOMAINS

AI & RAG Infrastructure	Data Engineering & Analytics	Hardware Acceleration
RAG Pipelines & Vector Databases	PySpark, DuckDB, PostgreSQL	CUDA GPU Edge Acceleration
PyTorch, TensorFlow, Keras, ONNX	Scientific Computing (M.Sc. Vienna)	Embedded Vision Pipelines
LLM Orchestration (LangChain, LangGraph)	Apache Superset & Data Pipelines	Low-Latency Inference Engines

CORE TECHNICAL SPECIALIZATION – ML & DATA SYSTEMS

- **Production AI Infrastructure:** Designing scalable RAG architectures, vector database integration, and high-throughput LLM inference backends.
- **Hardware Acceleration & Edge ML:** Optimizing deep learning models (ONNX, PyTorch) via CUDA for low-latency execution on GPU edge hardware.
- **Mathematical & Data Pipelines:** Applying numerical algorithms, distributed data processing (PySpark, DuckDB), and automated verification to data workflows.

Professional Experience

Mindbreeze GmbH – Software / Quality Engineer (C++/Python) 09/2022 – 07/2026
C++, Python, REST APIs, Linux, Docker, Azure DevOps, Telemetry & Performance Profiling | Team: 6 Engineers

- Optimized enterprise search index data pipelines for distributed data processing and large-scale document processing.
- Engineered Python automation frameworks to measure end-to-end data throughput and system response latencies.
- Resolved data pipeline memory degradations and performance bottlenecks during high-volume production releases.

Untis GmbH – Software Engineer (C++) 02/2022 – 06/2022
C++, System Refactoring, Performance Profiling | Team: 10 Engineers

- Refactored core modules to optimize memory allocation and algorithmic efficiency.

logi.cals – Embedded Software Engineer (C/C++) 02/2021 – 01/2022
C, C++, Embedded Systems, Deterministic Execution | Team: 6 Engineers

- Engineered deterministic C/C++ runtime execution systems and automated verification suites for complex hardware environments.

EYYES – Software Engineer (C++/Python) 04/2017 – 12/2020
C++, Python, Real-Time Execution Pipelines, CUDA, Embedded GPU Edge | Team: 8 Engineers

- Developed computer vision and deep learning pipelines (Keras, ONNX) for industrial camera systems to ensure continuous real-time edge performance.
- Engineered low-latency C++ and CUDA execution routines for real-time video stream processing on embedded GPU edge devices.
- Optimized detection accuracy and recognition precision through rigorous algorithm tuning and automated benchmarking suites.

Selected Software Projects

- **Azure AI Document Analyzer & RAG Infra:** RAG processing pipeline utilizing vector embeddings, LLM orchestration, and Azure cloud integration (GitHub).
  Tech: Python, RAG, Vector DBs, LLMs, Azure
- **AWS Cloud Deep Learning Pipeline:** Deep learning vision pipeline with automated data ingestion, model evaluation, and PyTorch inference (GitHub).
  Tech: PyTorch, Computer Vision, AWS, MLOps, Python
- **Magnum C++ Game of Life:** Parallel grid simulation exploring computational cellular automata and algorithmic performance (GitHub).
  Tech: Modern C++, Algorithmic Optimization, Parallel Computing

Technical Skills

ML, AI & Data: PyTorch, TensorFlow, Keras, ONNX, RAG Infrastructure, Vector DBs, LangChain, LangGraph, PySpark, DuckDB, PostgreSQL, Apache Superset.
Core Systems & GPU: Modern C++, Python, CUDA, OpenCL, OpenMP, MPI, Linux, Docker, REST APIs.
Cloud & DevOps: Amazon AWS, Microsoft Azure, GitLab CI/CD, Git, Telemetry, Cursor, GitHub Copilot.

Education

M.Sc. Scientific Computing – University of Vienna 10/2013 – 06/2017
Thesis: Silent Data Corruption in FT-GMRES (Link, Advisor: Prof. Dr. W. Gansterer).

B.Sc. Industrial Simulation – FH St. Pölten 09/2010 – 06/2013
Thesis: Simulation of Nanocomposite Materials (Link, Advisor: Prof. Dr. T. Schrefl).

Engineering Approach & Mindset

- **Systematic Problem Solving:** First-principles analytical approach to systematic debugging and stability optimization.
- **Quality & Efficiency Focus:** Pragmatic engineering mindset balancing rapid deployment with long-term maintainability and high performance.
- **Cross-Functional Integration:** Acts as the technical glue bridging development, QA, and architecture.

Languages & Work Status

German: Native (C2) · **English:** Full Professional Proficiency (C1/C2) · **Work Permit:** EU Citizen / Austria